Serious Shift — Architecture

How the pieces fit. One rule: Postgres is the single source of truth for all data.

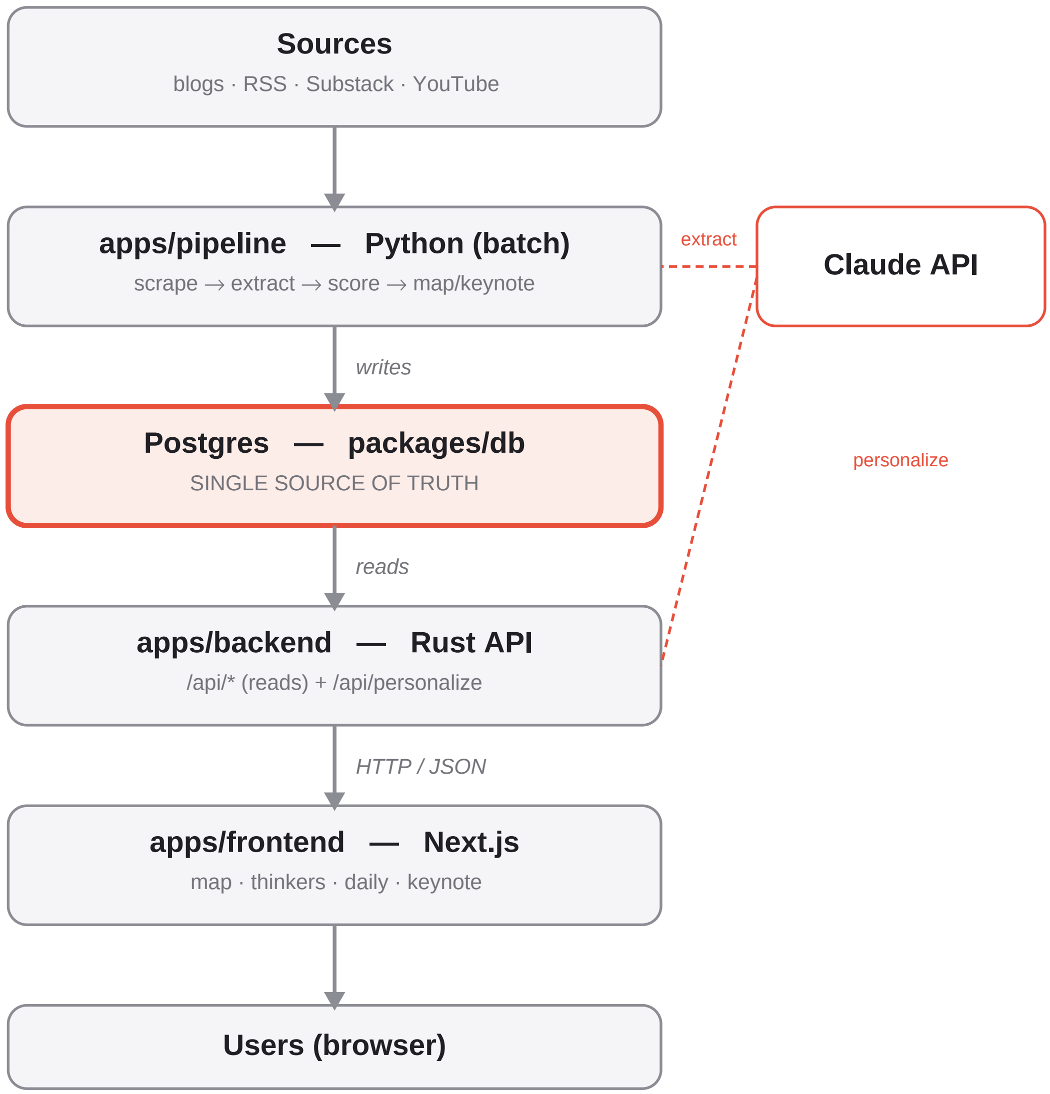

Data flows one way: pipeline writes the DB · backend reads it · frontend renders it. Each block deploys independently.

What each component does

apps/pipeline — Python, scheduled batch (the writer)

Scrapes each thinker's sources, extracts claims/predictions with Claude, scores them (recency + credibility), and generates the trend map & keynote. Run by run_weekly. Writes to Postgres.

Postgres — packages/db (the source of truth)

Holds everything: claims, predictions, the map, thinker bios/images, the scrape manifest, and the map/keynote/daily document blobs. Schema is versioned with dbmate migrations. No data lives in files.

apps/backend — Rust API (the reader)

Serves the data over HTTP. Each read endpoint is one SQL query that lets Postgres build the JSON, so it stays tiny. Also proxies Claude for /api/personalize. Replaces the old 53 MB of static JSON.

apps/frontend — Next.js app (the UI)

The trend map, thinker profiles, daily briefing, and keynote. Fetches everything from the backend API; no bundled data. Deploys free on Vercel.

Why it's built this way

- One source of truth (Postgres) — no data JSON, so nothing can drift from the database.
- Separate blocks, separate deploys — pipeline (batch) and backend (live API) have different runtimes.
- Python pipeline · Rust backend · Next.js frontend — each ecosystem where it is strongest.
- Backend lets Postgres assemble JSON — minimal, reviewable, matches the old data shapes exactly.
- dbmate migrations — language-neutral, so the schema is not tied to Python or Rust.

Full detail: README.md + ARCHITECTURE.md and each block's README. Table catalog: packages/db/README.md.